

Concurrency, Parallelism and Distributed Systems (CPDS)

Understanding parallelism

Eduard Ayguadé, Josep-Ramon Herrero, Daniel Jiménez
({eduard,josepr,djimenez}@ac.upc.edu)

Computer Architecture Department
Universitat Politècnica de Catalunya

2014/15 - Fall

Part I

A brief introduction to parallel architectures

Outline

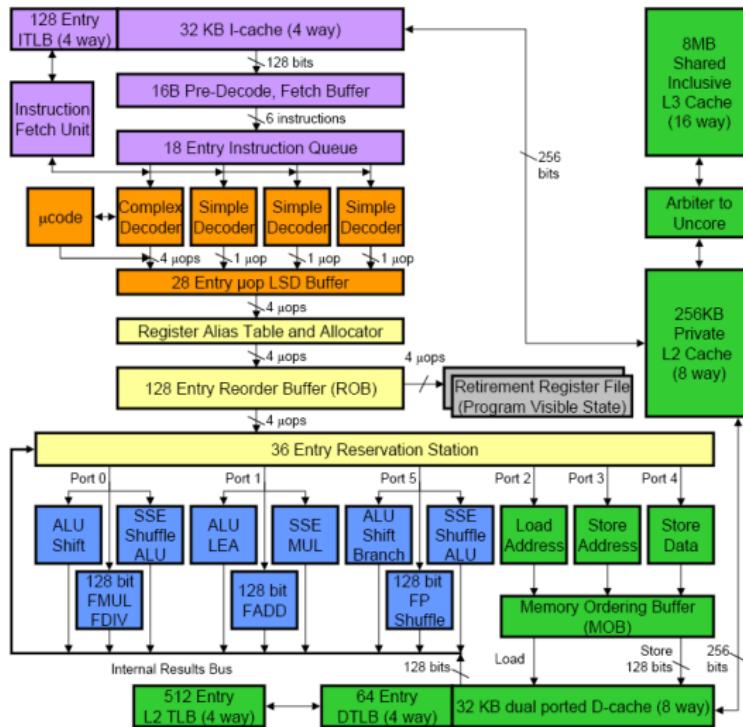
Uniprocessor parallelism

Shared-memory architectures

Distributed-memory architectures

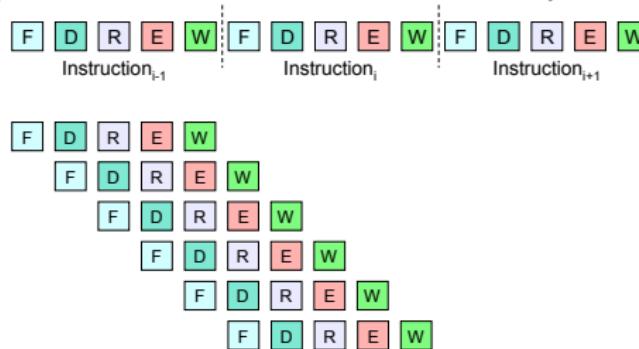
Graphic Processing Unit (GPU)

Current uniprocessor architecture: Intel Nehalem i7



Pipelining

- ▶ Execution of single instruction divided in multiple stages
- ▶ Overlap the execution of different stages of consecutive instructions
- ▶ Ideal: $IPC=1$ (1 instruction executed per cycle)



- ▶ $IPC < 1$ due to hazards (structural, data, control), preventing the execution of an instruction in its designated clock cycle

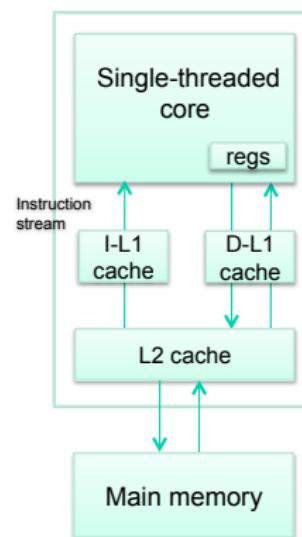
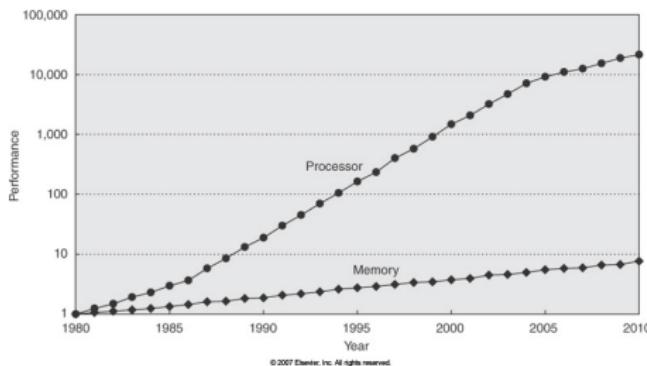
Sources of parallelism in uniprocessors

- ▶ ILP (Instruction-level parallelism)
 - ▶ Superscalar architecture: multiple issue slots (functional units)
 - ▶ Execution of multiple instructions, from the same instruction flow, per cycle
- ▶ TLP (thread-level parallelism)
 - ▶ Multithreaded architecture¹: fill the pipeline with instructions from multiple instruction flows
 - ▶ Latency hiding (cache misses, non-pipelined FP, ...)
- ▶ DLP (data-level parallelism)
 - ▶ SIMD architecture: single-instruction executed on multiple-data in a single word
 - ▶ Vector functional unit

¹Hyperthreading in Intel terminology

Memory hierarchy

- ▶ Addressing the yearly increasing gap between CPU cycle and memory access times



- ▶ Size vs. access time

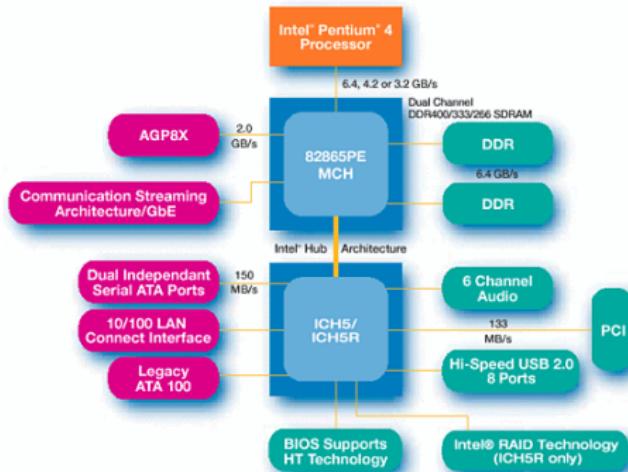
- ▶ Non-blocking design

Memory hierarchy

- ▶ The principle of locality: if an item is referenced ...
 - ▶ Temporal locality: ... it will tend to be referenced again soon (e.g., loops, reuse)
 - ▶ Spatial locality: ... items whose addresses are close by tend to be referenced soon (e.g., straight line code, array access)
- ▶ Line (or block)
 - ▶ A number of consecutive words in memory (e.g. 32 bytes, equivalent to 4 words x 8 bytes)
 - ▶ Unit of information that is transferred between two levels in the hierarchy
- ▶ On an access to a level in the hierarchy
 - ▶ Hit: data appears in one of the lines in that level
 - ▶ Miss: data needs to be retrieved from a line in the next level

Components in a node board

- ▶ The processor is connected to memory (and graphics unit or GPU) via the so-called 'northbridge'
- ▶ Through the northbridge, the processor is connected to slower I/O devices through the so-called 'southbridge'

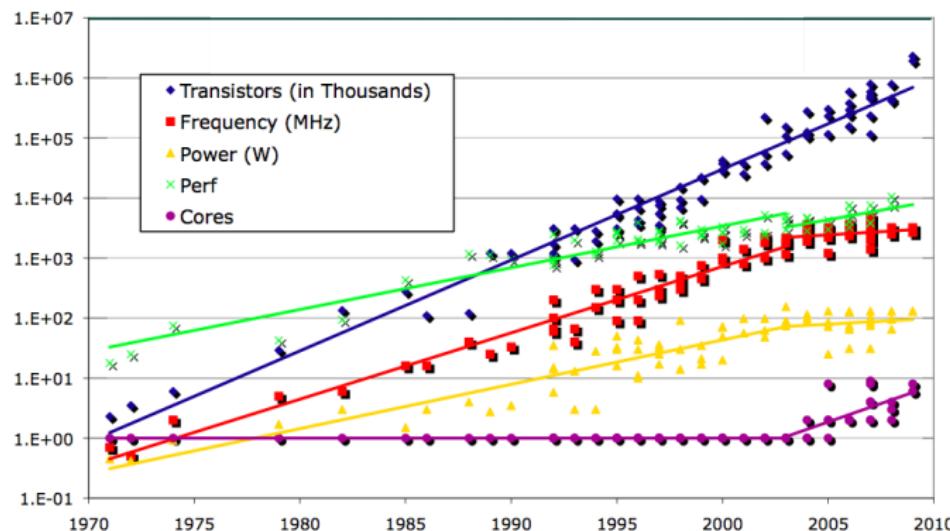


Who exploits this uniprocessor parallelism?

- ▶ In theory, the compiler understands all of this ... but in practice the compiler may need your help:
 - ▶ Software pipelining to statically schedule ILP
 - ▶ Unrolling to allow the processor dynamically exploits ILP
 - ▶ Data contiguous in memory and aligned to efficiently exploit DLP
 - ▶ Blocking (or tiling) to define a problem that fits in register/L1-cache/L2-cache
 - ▶ ...
- ▶ Optimized libraries usually help for most commonly used operations

Transistors, frequency, power, performance and ... cores!

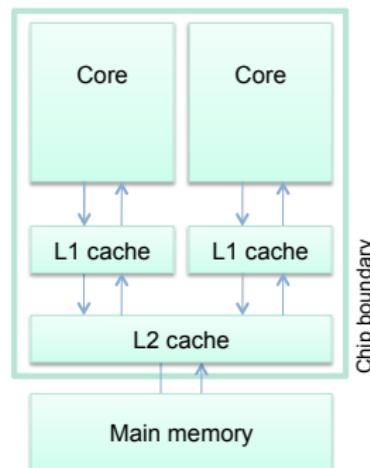
An inflection point in 2004 ... the power wall².



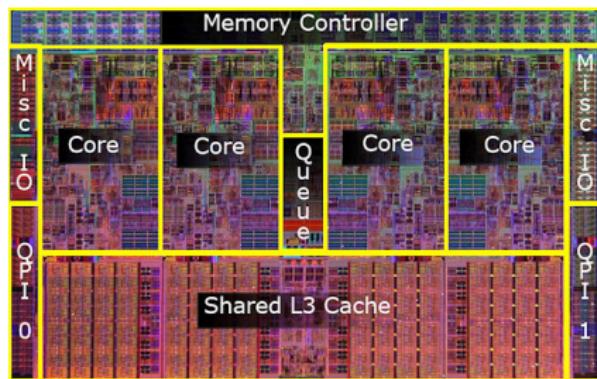
² Data from K. Olukotun, L. Hammond, H. Sutter, B. Smith, C. Batten, and K. Asanovic.

Multicore era

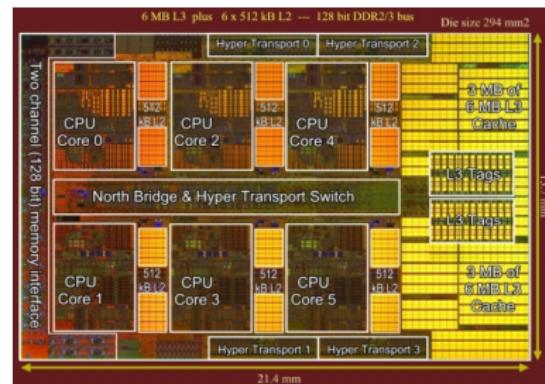
- ▶ The number of transistors on a chip is continuing to increase to accommodate multiple processors (cores) on a single chip
- ▶ Multicore = Chip Multi-Processor (CMP)



Multicore era (examples)

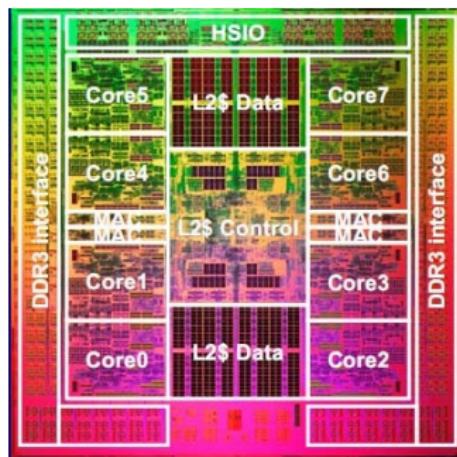


Intel's Nehalem i7 4-core (2 threads/core), 32/32 KB L1 and 256 KB L2 per core. Unified 4 or 8 MB L3

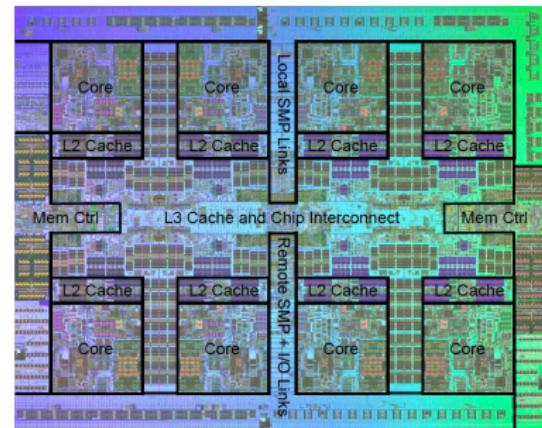


AMD's Istanbul 6-core, 512KB L2 per core and 6 MB shared L3

Multicore era (examples)



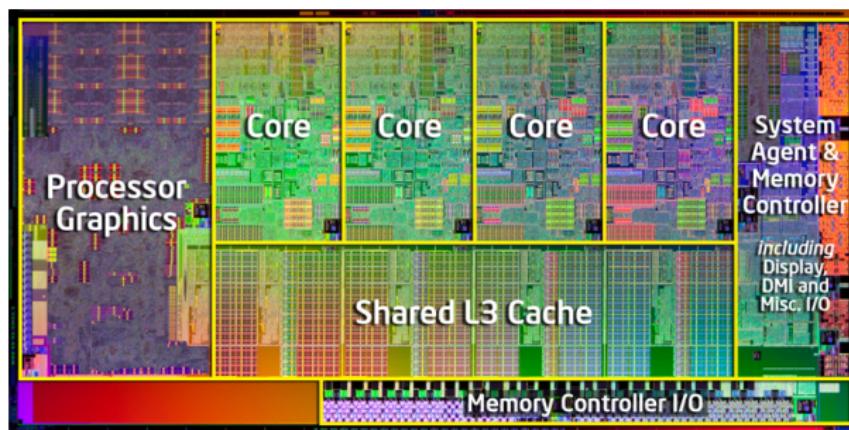
Sun's SPARC64-VIIIfx 8-core, 32/32K L1 per core and 5 MB shared L2. 2 GHz.. 760 M transistors



IBM's Power7 8-core (4 threads/core), 32/32K L1 and 256KB L2 per core. 32 MB shared L3. 2.4 GHz. 1.2 billion transistors

Multicore era (examples)

- ▶ Homogeneous multi-core systems include only identical cores, heterogeneous multi-core systems have cores with different functionalities, performances, ... (e.g. Intel SandyBridge)



Outline

Uniprocessor parallelism

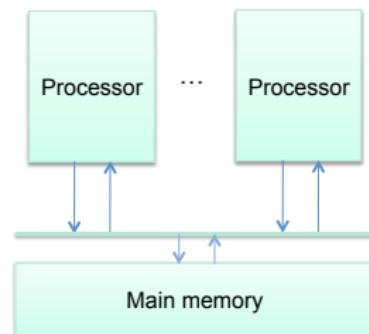
Shared-memory architectures

Distributed-memory architectures

Graphic Processing Unit (GPU)

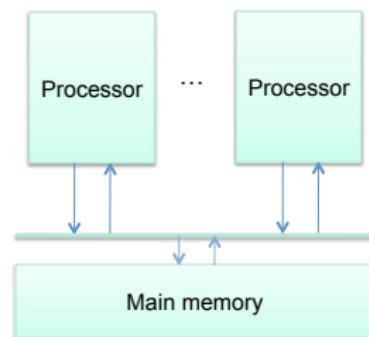
Shared-memory multiprocessors

- ▶ SMP (Symmetric Multi-Processor)
 - ▶ Two or more identical processors are connected to a single shared main memory
 - ▶ Interconnection network: any processor can access to any memory location
- ▶ Symmetric multiprocessing: a single OS instance on the SMP
- ▶ Asymmetric multiprocessor (e.g. high/low ILP processors, ...) and multiprocessing (e.g. some processors running OS, others user code)



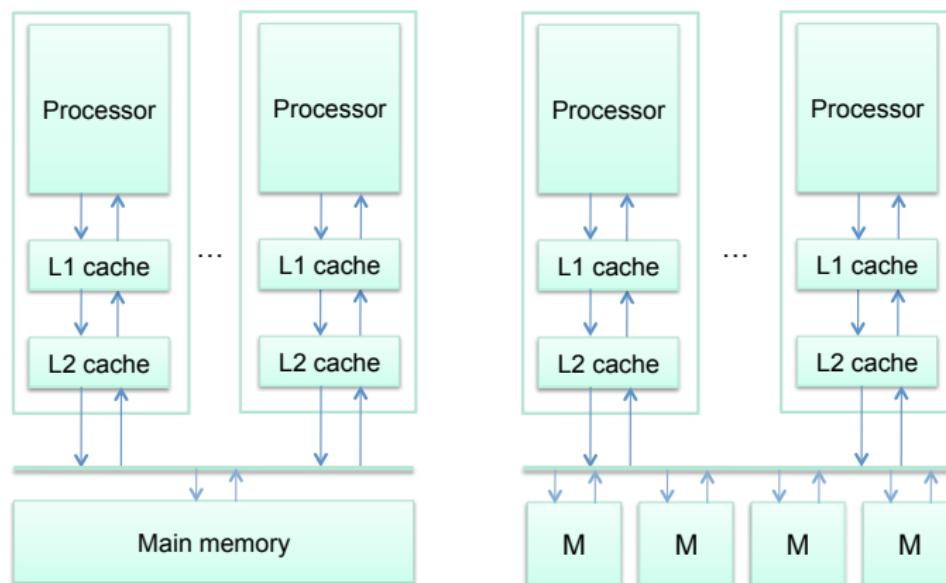
Shared-memory multiprocessors

- ▶ SMP (Symmetric Multi-Processor)
 - ▶ Access to shared data with load/store instructions
 - ▶ Uniform Memory Access (UMA): access time to a memory location is independent of which processor makes the request or which memory chip contains the data
- ▶ The bottleneck in the scalability of SMP is the 'bandwidth' of the interconnection network and the memory



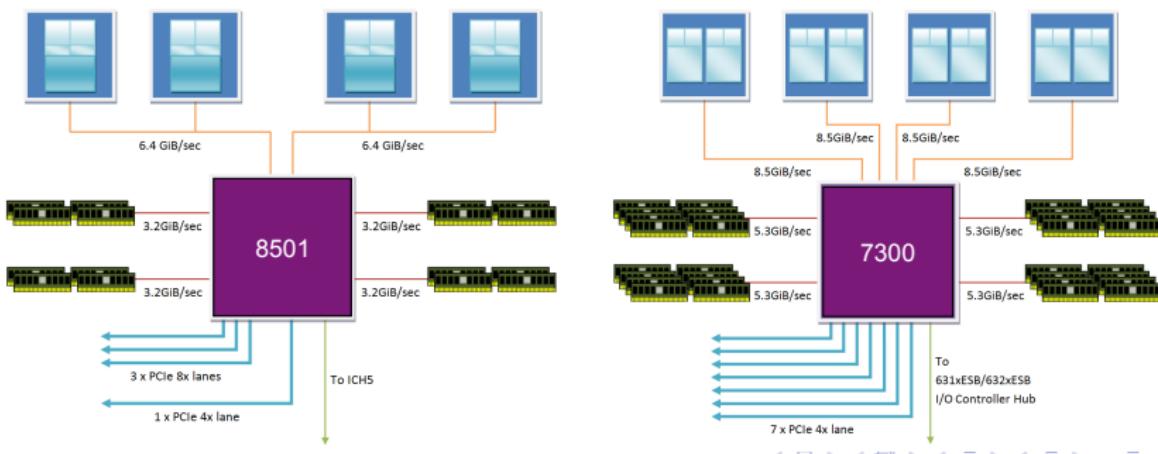
Shared-memory multiprocessors

Local caches and multi-banked (interleaved) memory

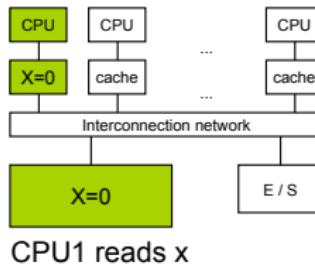


Front-Side Bus (FSB)

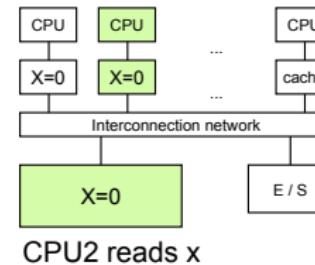
- ▶ Processor-northbridge connection (64-bit, 1-4 transfers/cycle)
- ▶ A number of processors can be connected to a single FSB sharing its total bandwidth
- ▶ Northbridge with several FSB ports



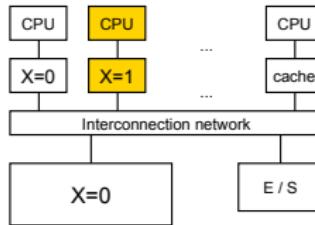
Shared-memory multiprocessors: coherence



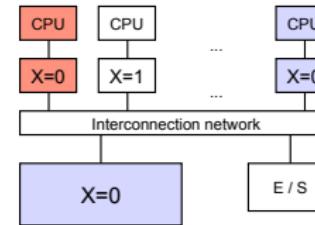
CPU1 reads x



CPU2 reads x



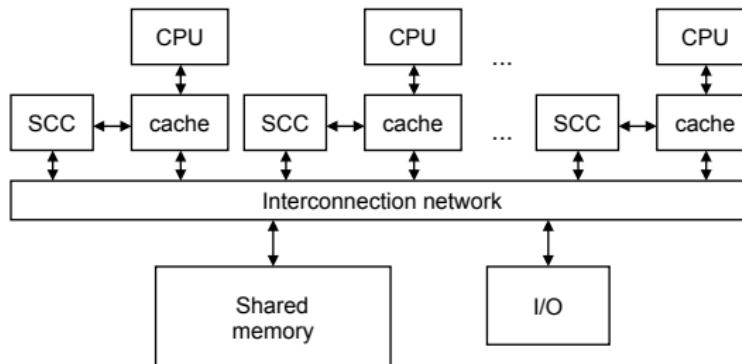
CPU2 writes x

CPU1 or CPU3
read an incorrect x

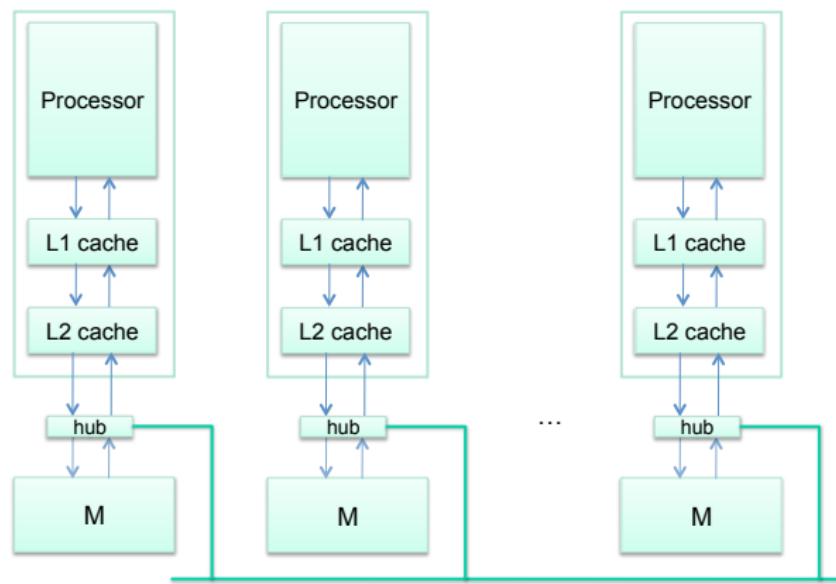
Snooping-based coherence protocol

Shared bus is a broadcast medium

- ▶ Transactions on bus are visible to all processors
- ▶ Processors or their representatives can snoop (monitor) bus and take action on relevant events (e.g. change state)



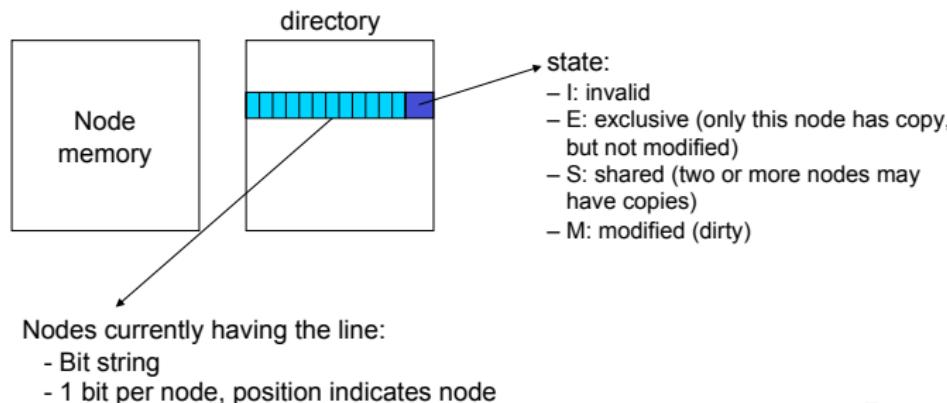
Escalability? Non-Uniform Memory Architectures



- ▶ Hub enables cache-coherent NUMA

Directory-based cache coherency protocol

- ▶ Directory-based is more scalable than Snoopy (bus)
- ▶ Only processors with copies receive info \Rightarrow Reduce Communication
 - ▶ It tracks the state of the data stored in its memory
 - ▶ If a block is shared, it must track which processors have it
 - ▶ If a block is valid only in one node, it should know which one



Shared-memory multiprocessors: synchronization

- ▶ Need hardware support to guarantee atomic (indivisible) instruction to fetch and update memory
 - ▶ User-level synchronization operations (e.g. locks, barriers, point-to-point, ...) using these primitives
- ▶ **test-and-set**: read value in location and set to 1

```
lock: t&s r2, flag  
      bnez r2, lock    // already locked?
```

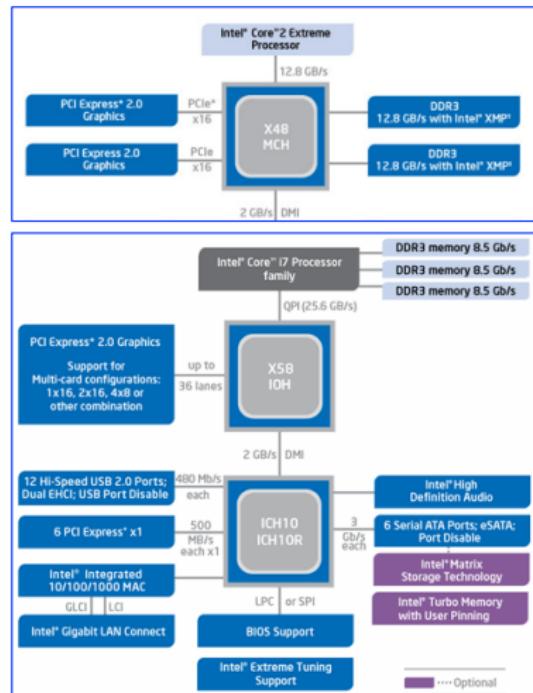
- ▶ Atomic exchange: interchange of a value in a register with a value in memory
- ▶ **fetch-and-op**: read value in location and replace with increment/decrement

Shared-memory multiprocessors: synchronization

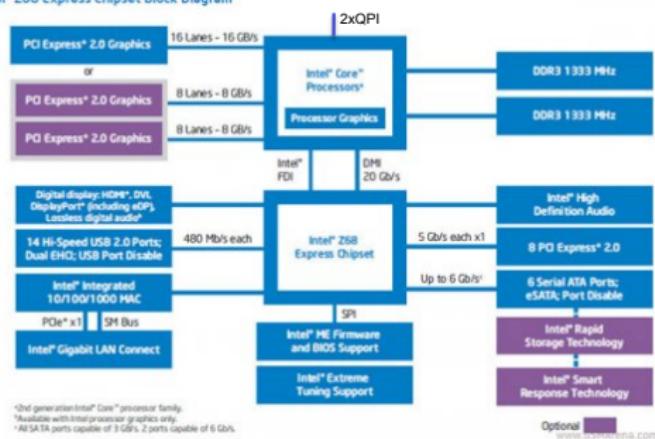
- ▶ Atomicity difficult or inefficient in large systems. Idea: assume optimistic atomic access
- ▶ Load-linked Store-conditional ll-sc
 - ▶ ll returns the current value of a memory location
 - ▶ sc stores a new value in that memory location if no updates have occurred to it since the ll; otherwise, the sc is guaranteed to fail

```
lock: ll r1, location
      bnez r1, lock
      r2 = f (r1)           // f can be test-and-set, fetch-and-op, ...
      sc location, r2
      beqz lock            // Z flag set to 1 if atomicity has occurred
```

Northbridge evolution



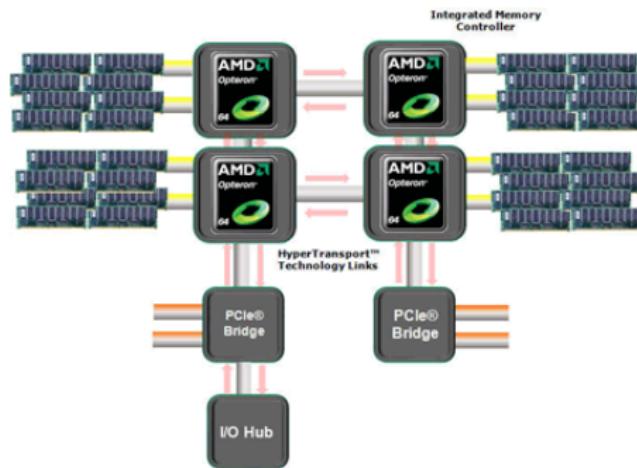
Intel® Z68 Express Chipset Block Diagram



- Memory controller inside the processor chip (e.g. Intel Nehalem)
- Graphics unit inside the processor chip (e.g. Intel Sandybridge)

Quick Path Interconnect (QPI) and Hypertransport (HT)

- ▶ Point-to-point processor interconnects, replacing the FSB
- ▶ Processor chips with integrated memory controllers
- ▶ Connect one or more processors and/or one or more chipsets in a network: non-uniform memory access (NUMA)



Outline

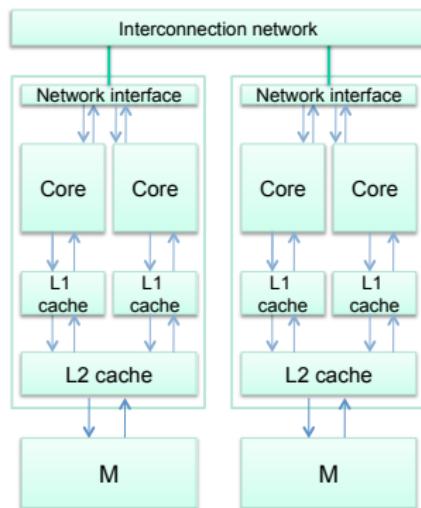
Uniprocessor parallelism

Shared-memory architectures

Distributed-memory architectures

Graphic Processing Unit (GPU)

Why hardware needs to provide data sharing?



- ▶ Simple design: distributed-memory architectures
- ▶ Each node has its own (local) exclusive memory, accessible through load/store instructions (no cache coherency among nodes)
- ▶ Network interface and routers to exchange data between nodes (interconnection network)

Interconnection networks

- ▶ Interconnection networks are build up of switching elements
 - ▶ Switches: devices that contain multiple input and output ports with a crossbar interconnection between them (i.e. any input to any output path available)
- ▶ Topology is the pattern in which the individual switches are connected to other switches and to processors and memories (nodes).
 - ▶ Direct topologies connect each switch directly to the network interface of a node
 - ▶ In indirect topologies at least some of the switches connect to other switches
- ▶ Influence on latency, bandwidth and congestion

Interconnection networks: direct topologies

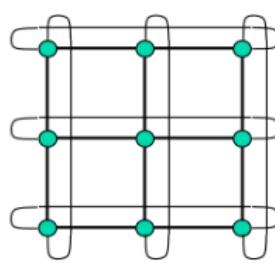
line



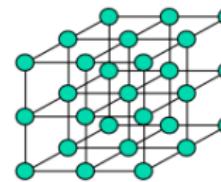
ring (1D torus)



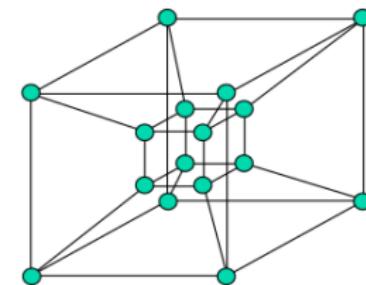
2D torus



3D mesh

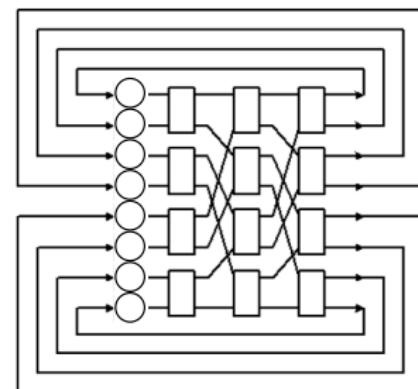
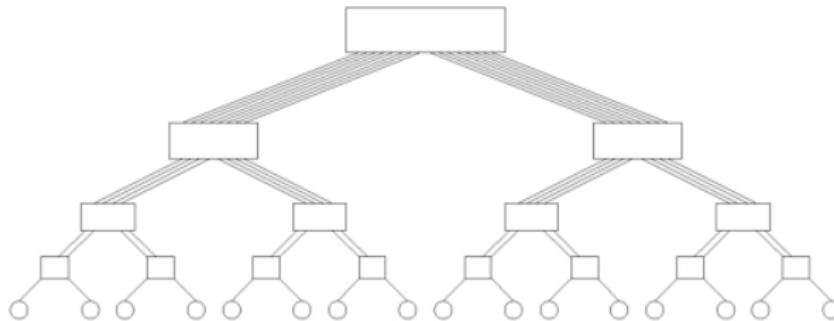


hipercube



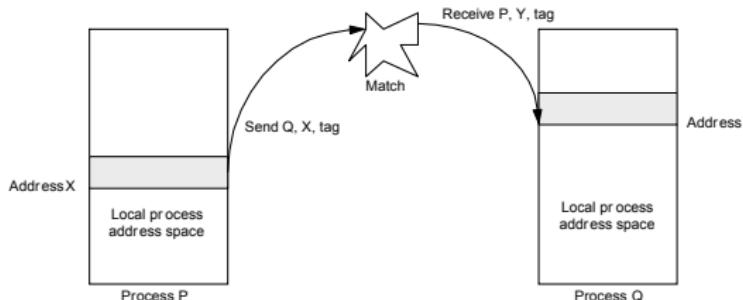
Interconnection networks: indirect topologies

E.g. fat tree (left) and Omega multistage (right) networks



Communication model

Data exchange using send and receive primitives



- ▶ Send specifies buffer to be sent and receiving process
- ▶ Receive specifies sending process and application storage to receive into
- ▶ Optional tag on send and matching rule on receive
- ▶ Optional implicit synchronization (e.g. blocking receive)

Who does communication?

- ▶ Software DSM (distributed-shared memory)
 - ▶ Software layer that implements data sharing (and coherence)
 - ▶ Transparently to programmer
 - ▶ Usually based on page faults (OS involved, high overhead), which uses the communication model to move pages between nodes
- ▶ Message-passing paradigm
 - ▶ User-level library exporting the communication model to the programmer
 - ▶ Programmer moves data when necessary, assuming a data distribution

Outline

Uniprocessor parallelism

Shared-memory architectures

Distributed-memory architectures

Graphic Processing Unit (GPU)

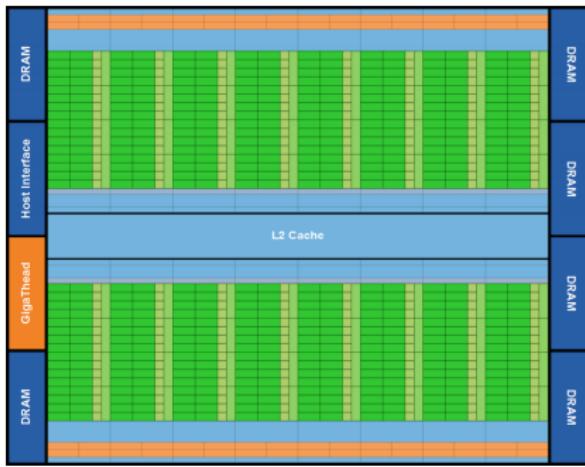
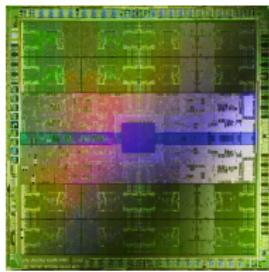
Why GPU architectures

GPU are specialized for highly parallel, compute-intensive computation

- ▶ Simple control: same computation on all compute units
- ▶ Massive number of threads to reduce impact of long latency memory accesses



NVIDIA GPU architecture

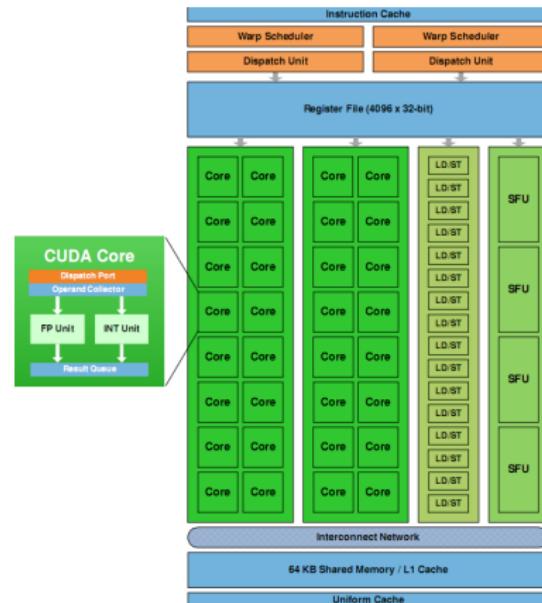


Die and aerial view of the NVIDIA Fermi architecture, with the graphics-specific parts (raster engine, polymorphic engines, ...) omitted.

NVIDIA GPU architecture (cont.)

The NVIDIA Fermi architecture consists of:

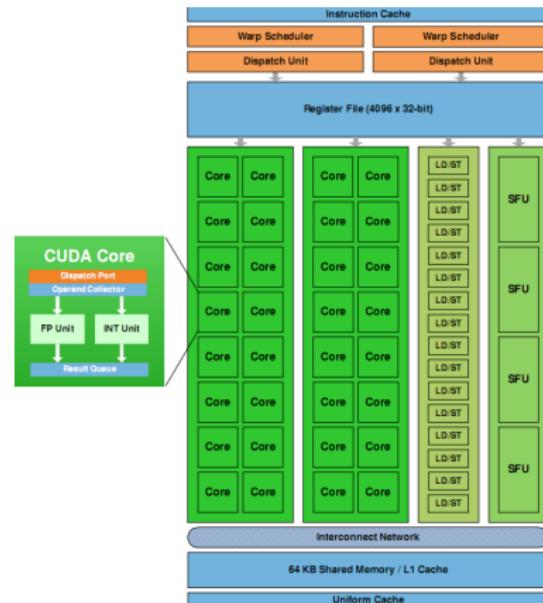
- ▶ 16 streaming multiprocessors (SM) interconnected via a 768k L2 cache crossbar
 - ▶ Each SM has 32 cores and 4096 registers, sharing 64 KB (48K L1 and 16K shared memory or viceversa).
 - ▶ Each core has one FP unit and one integer unit



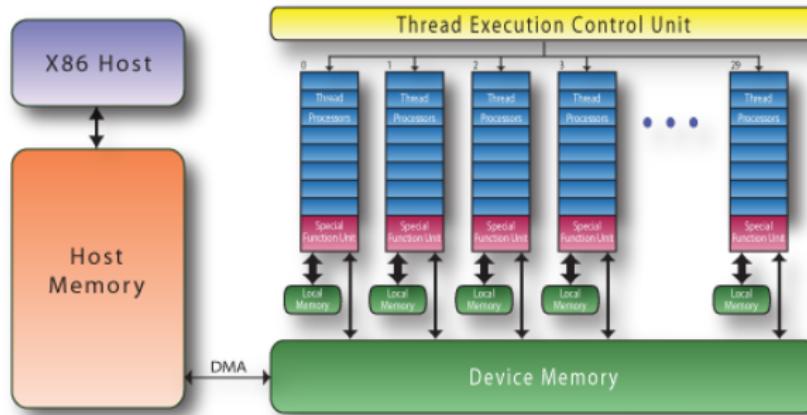
NVIDIA GPU architecture (cont.)

The NVIDIA Fermi architecture consists of:

- ▶ Each SM executes a *warp* of 32 threads, can multiplex 48 active warps
- ▶ Each SM has 16 load/store units, allowing source and destination addresses to be calculated for sixteen threads per clock
- ▶ Each SM also has 4 Special Function Units (SFU) that execute complex instructions such as sin, cosine, reciprocal, and square root



CPU-GPU block diagram



Separate address spaces between CPU (host memory) and GPU (device memory). Communication using IO commands and DMA memory transfers (e.g. PCI Express)

Part II

Parallelism: concept and metrics

Outline

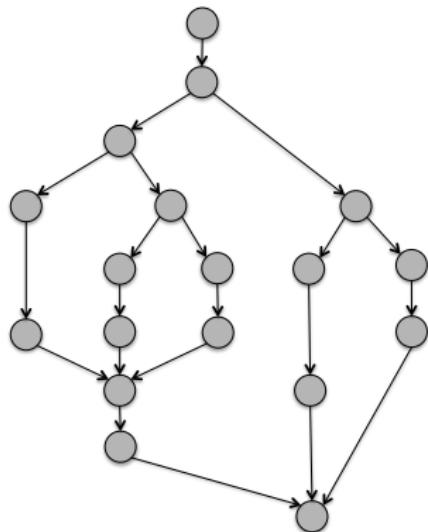
Parallelism

Speedup and Amdahl's law

Finding concurrency/parallelism

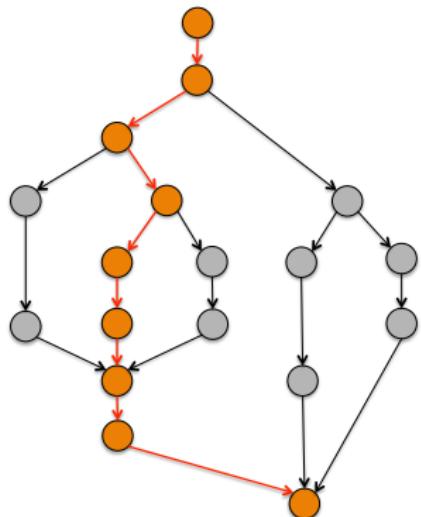
- ▶ Can the computation be divided in parts?
 - ▶ Based on the processing to do:
 - ▶ Task decomposition (e.g. functions, loop iterations)
 - ▶ Based on the data to be processed:
 - ▶ Data decomposition (e.g. matrix) (implies task decomposition)
- ▶ There may be data dependencies between tasks
- ▶ The decomposition determines the potential parallelism that could be obtained

Computing the parallelism



- ▶ Computation task graph abstraction
 - ▶ Directed Acyclic Graph
 - ▶ Node = arbitrary sequential computation
 - ▶ Edge = dependence (successor node can only execute after predecessor node has completed: dataflow)
- ▶ Processor abstraction (simplification)
 - ▶ P identical processors
 - ▶ Each processor executes one node at a time

Computing the parallelism



- ▶ $T_1 = \sum_{i=1}^{nodes} (work_node_i)$
- ▶ $T_\infty = \sum_{i \in criticalpath} (work_node_i)$, assuming sufficient resources
- ▶ $Parallelism = T_1/T_\infty$, independent of number of processors P
- ▶ P_{min} is the minimum number of processors necessary to achieve *Parallelism*

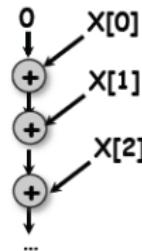
Example: vector sum

Compute the sum of elements $X[0] \dots X[n-1]$ of a vector X

Sequential algorithm.

```
sum = 0; for ( i=0 ; i< n ; i++ ) sum += X[i];
```

- ▶ Computation graph



$$T_1 = O(n)$$

$$T_\infty = O(n)$$

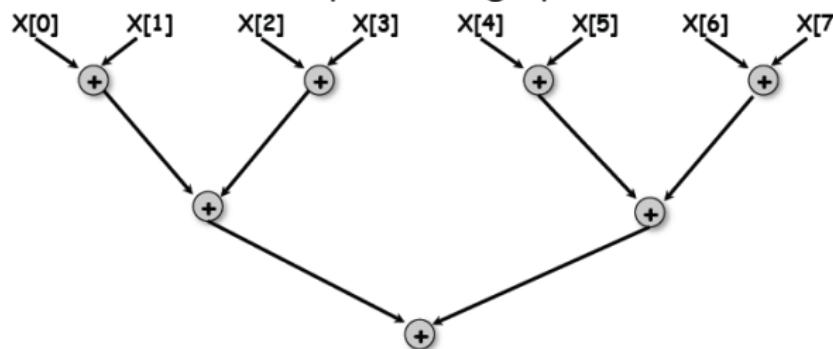
$$\text{Parallelism} = O(1)$$

How can we design an algorithm (computation graph) with more parallelism?

Example: vector sum

Compute the sum of elements $X[0] \dots X[n-1]$ of a vector X

- An alternative computation graph



- $T_1 = O(n); T_\infty = O(\log n); \text{Parallelism} = O(n/(\log n))$
- How to restructure the sequential algorithm to have this computation graph? (recursive solutions)

Example: vector sum

Compute the sum of elements $X[0] \dots X[n-1]$ of a vector X

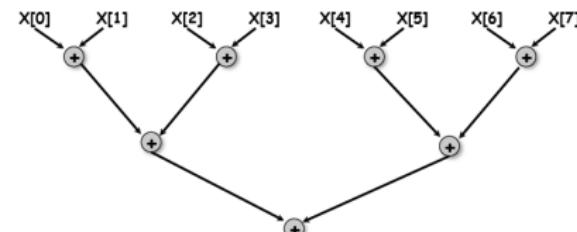
Recursive algorithm:

```
int recursive_sum(int *X, int n)
{
    int ndiv2 = n/2;
    int sum=0;

    if (n==1) return X[0];

    sum = recursive_sum(X, ndiv2);
    sum += recursive_sum(X+ndiv2, n-ndiv2);
    return sum;
}

void main()
{
    int sum, X[N];
    ...
    sum = recursive_sum(X,N);
    ...
}
```

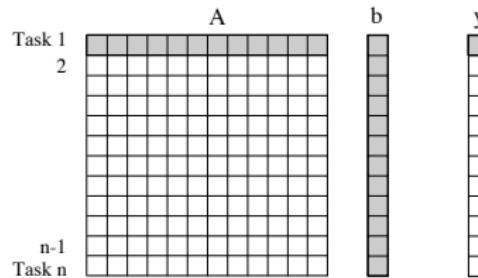


Granularity and parallelism

- ▶ Each node in the computation task graph represents a sequential computation
- ▶ The granularity of the decomposition is determined by the size of each node in the computation graph
- ▶ Fine-grained tasks vs. coarse-grained tasks: the degree of parallelism increases as the decomposition becomes finer in granularity and vice versa

Granularity and parallelism: fine-grained decomposition

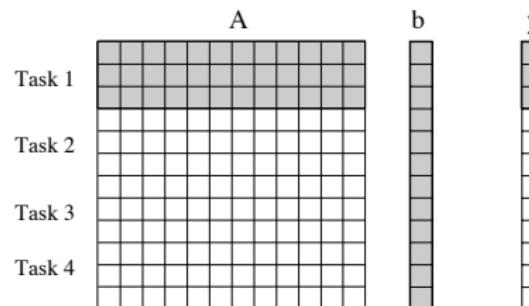
Example: matrix-vector product (n by n matrix):



- ▶ A task could be computing the whole vector y
- ▶ A task could be each individual \times or $+$ in the dot product that computes an element of y ($y[i] = y[i] + A[i][j] * b[j]$)
- ▶ A task could also be each complete dot product to compute an element of y ($y[i] = \sum_{j=1}^{j=n} (A[i][j] * b[j])$)

Granularity and parallelism: coarse-grained decomposition

A coarse grained counterpart to the matrix-vector product example. Each task in this example corresponds to the computation of three elements of the result vector.



Granularity and parallelism: fine vs. coarse-grained

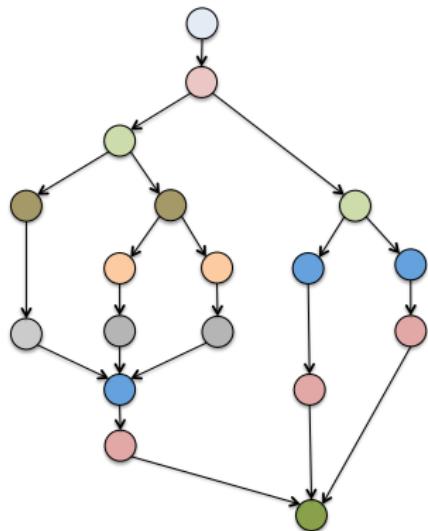
- ▶ It would appear that the parallel time can be made arbitrarily small by making the decomposition finer in granularity but...
 - ▶ Inherent bound on how fine the granularity of a computation can be
 - ▶ e.g. *matrix-vector multiply*: (n^2) concurrent tasks.
 - ▶ Tradeoff between the granularity of a decomposition and associated overheads (sources of overhead commented later, e.g. creation of tasks, synchronization, exchange of data between tasks, ...)
 - ▶ The granularity may determine performance bounds

Outline

Parallelism

Speedup and Amdahl's law

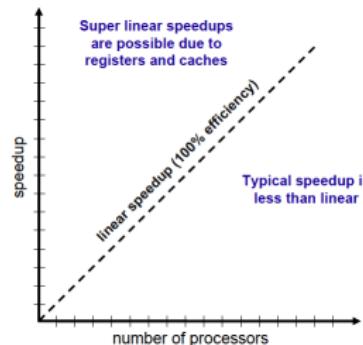
Speedup



- ▶ T_p = execution time on P processors
(depends on the schedule of the graph nodes on the processors)
- ▶ Lower bounds
 - ▶ $T_p \geq T_1/P$
 - ▶ $T_p \geq T_\infty$
- ▶ Speedup on P processors: $S_p = T_1/T_p$

Speedup vs. efficiency

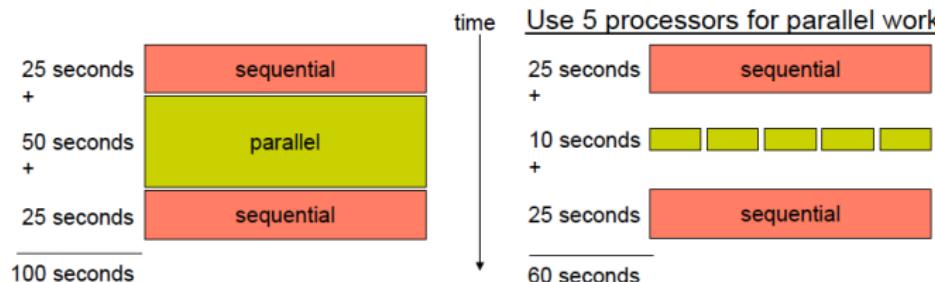
- ▶ Speedup S_p : relative reduction of execution time when using P processors with respect to sequential



- ▶ Efficiency Eff_p : it is a measure of the fraction of time for which a processing element is usefully employed
 - ▶ $Eff_p = T_1 / (T_p \times P)$
 - ▶ Also, $Eff_p = S_p / P$

Amdahl's law

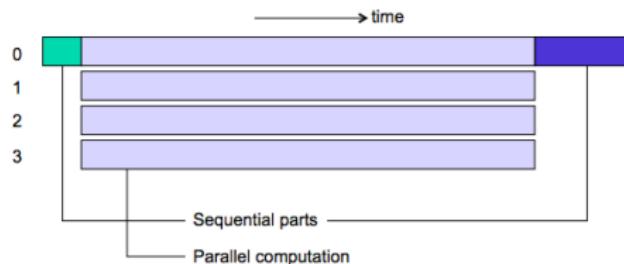
The performance improvement to be gained from using some faster mode of execution is limited by the fraction of the time the faster mode can be used



- ▶ Parallel part is 5 times faster: $Speedup_{parallel_part} = 50/10 = 5$
- ▶ Parallel version is just 1.67 times faster: $S_p = 100/60 = 1.67$

Amdahl's law

- ▶ Performance improvement is limited by the fraction of time the program runs in parallel (e.g. φ)



$$T_1 = T_{seq} + T_{par} = (1 - \varphi) \times T_1 + \varphi \times T_1$$

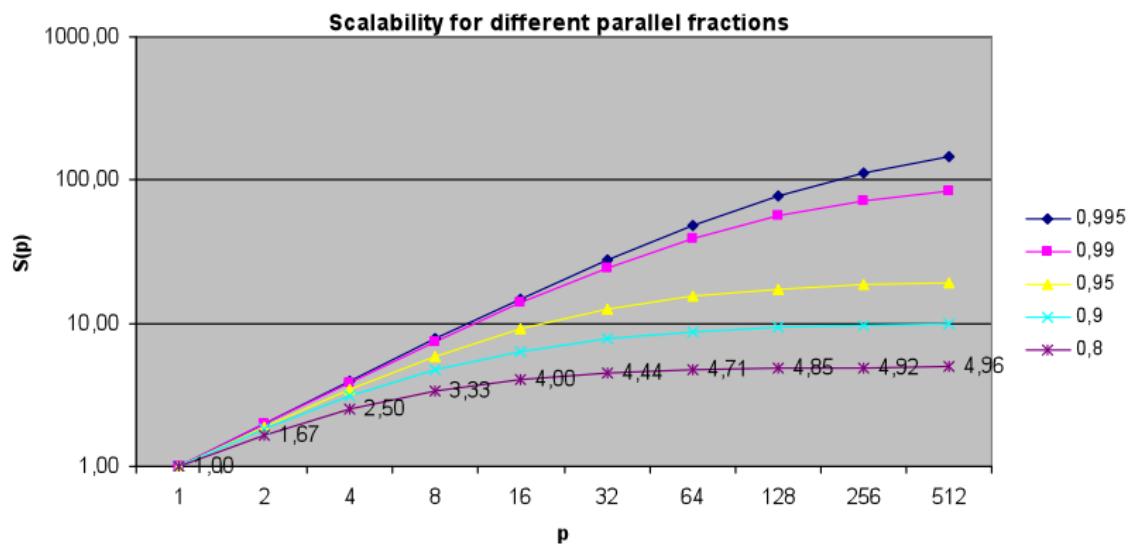
$$T_p = (1 - \varphi) \times T_1 + (\varphi \times T_1 / P)$$

$$S_p = \frac{T_1}{T_p}$$

$$S_p = \frac{1}{((1 - \varphi) + \varphi/P)}$$

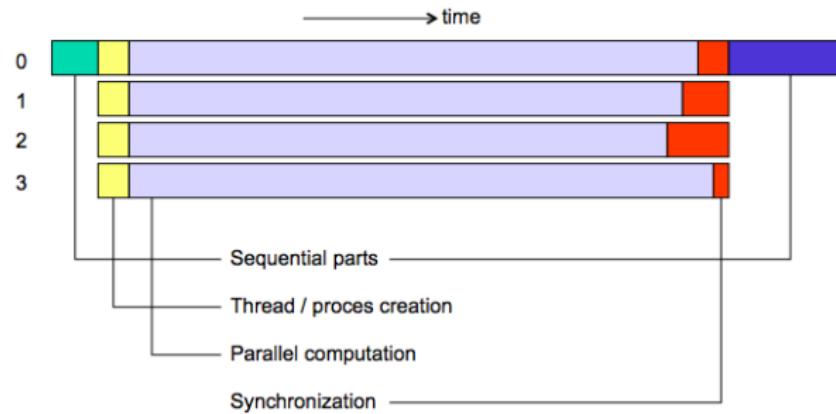
$$S_p \rightarrow \frac{1}{(1 - \varphi)} \text{ when } P \rightarrow \infty$$

Amdahl's law



Sources of overhead

Parallel computing is not free, we should account overheads (i.e. any cost that gets added to a sequential computation so as to enable it to run in parallel)



Sources of overhead

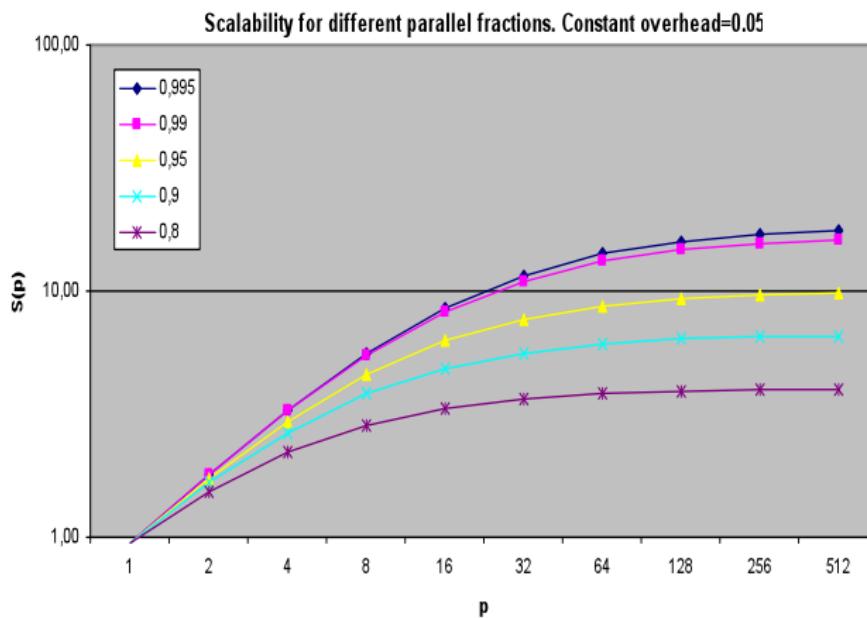
- ▶ **Task creation and termination:** extra processing performed at the start and end of each task
- ▶ **Synchronization:** extra processing to ensure that dependences in computation graph are satisfied
- ▶ **Data sharing:** can be explicit via messages, or implicit via a memory hierarchy (caches)

Sources of overhead (cont.)

- ▶ **Idleness:** thread cannot find any useful work to execute (e.g. dependences, load imbalance, poor communication and computation overlap or hiding of memory latencies, ...)
- ▶ **Computation:** extra work added to obtain a parallel algorithm (e.g. replication)
- ▶ **Memory:** extra memory used to obtain a parallel algorithm (e.g. impact on memory hierarchy, ...)
- ▶ **Contention:** competition for the access to shared resources (e.g. memory, network)

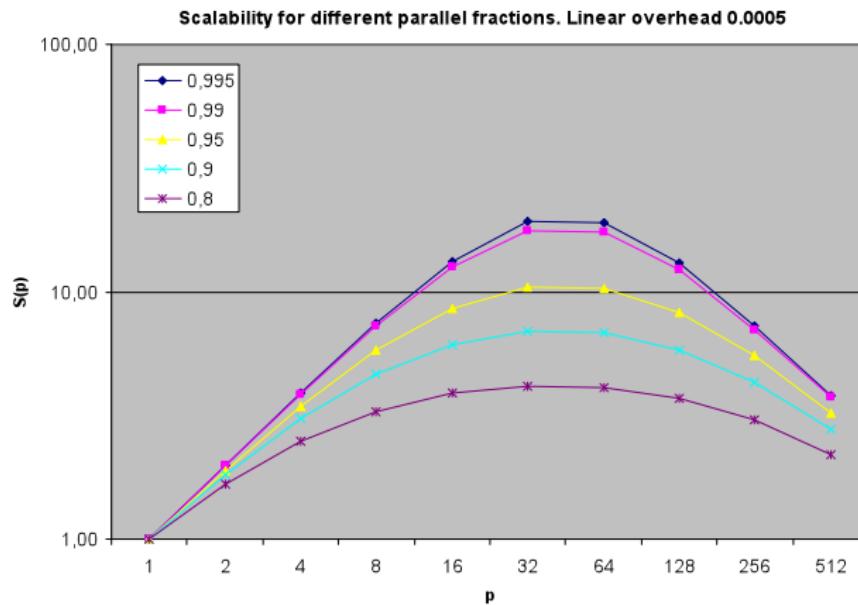
Amdahl's law (constant overhead)

$$T_p = (1 - \varphi) \times T_1 + \varphi \times T_1/p + \text{overhead}$$



Amdahl's law (linear overhead)

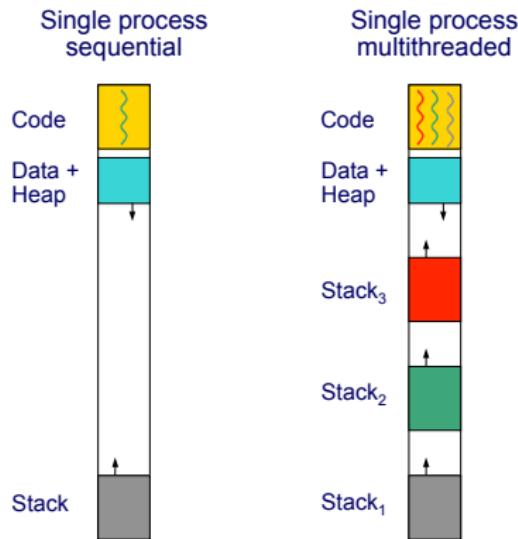
$$T_p = (1 - \varphi) \times T_1 + \varphi \times T_1/p + \text{overhead}_p$$



Part III

Address spaces in parallel programming

Shared memory: address space

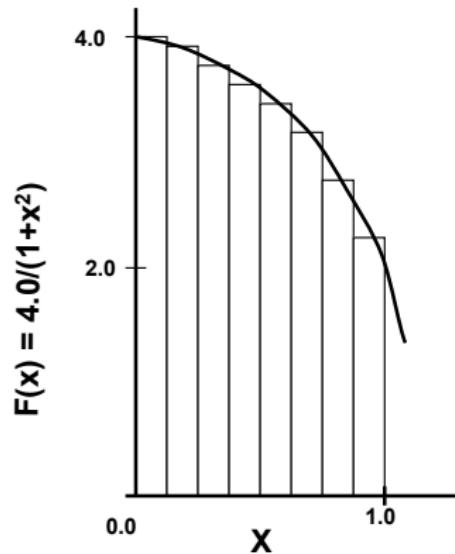


Programmer needs

- ▶ Distribute work
- ▶ All threads can access data, heap and stacks
- ▶ Use synchronization mechanisms to avoid data races

Shared memory programming: example

Mathematically, we know that:



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

Shared memory programming: example (cont.)

Sequential code:

```
void main ()  
{  
    int i;  
    double x, pi, step, sum = 0.0;  
    long num_steps = 100000;  
  
    step = 1.0/(double) num_steps;  
  
    for (i=1;i<= num_steps; i++) {  
        x = (i-0.5)*step;  
        sum = sum + 4.0/(1.0+x*x);  
    }  
    pi = step * sum;  
}
```

Work distribution: each processor responsible for computing some rectangles (in other words, to execute some iterations of loop *i*). Guarantee that all processors have computed their *sum* before combining into final value

Shared memory programming: example (cont.)

OpenMP code (not the shortest one, just for illustration purposes):

```
#include "omp.h"
#define NUM_THREADS = 4
void main ()
{
    int i, id;
    double x, pi, step, sum[NUM_THREADS];
    long num_steps = 100000;

    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);

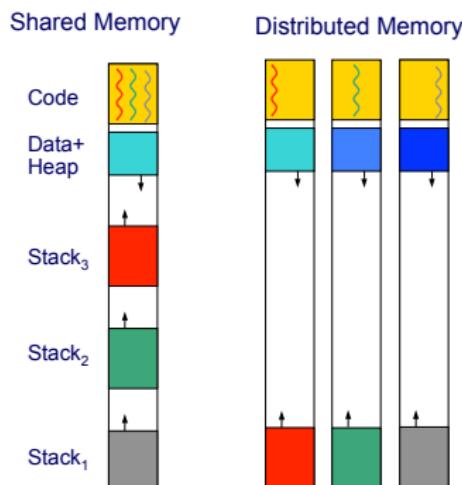
#pragma omp parallel private(id)
{
    id = omp_get_thread_num();
    sum[id] = 0.0;
#pragma omp for private(x,i)           // Implicit barrier synchronization at the end of for
    for (i=1;i<= num_steps; i++) {
        x = (i-0.5)*step;
        sum[id] += 4.0/(1.0+x*x);
    }
#pragma omp single                   // Only one computing final result
    for (i=0, pi=0.0; i<NUM_THREADS; i++)
        pi += step * sum[i];
}
}
```

Shared memory programming: example (cont.)

OpenMP code:

```
void main ()  
{  
    int i;  
    double x, pi, step, sum=0.0;  
    long num_steps = 100000;  
  
    step = 1.0/(double) num_steps;  
  
#pragma omp parallel for private(x) reduction(+:sum)  
    for (i=1;i<= num_steps; i++) {  
        x = (i-0.5)*step;  
        sum += 4.0/(1.0+x*x);  
    }  
    pi = step * sum;  
}
```

Distributed memory: address space



Programmer needs

- ▶ Distribute work
- ▶ Distribute data
- ▶ Use communication mechanisms to share data explicitly
- ▶ Use synchronization mechanisms to avoid data races

Distributed memory programming: example

Sequential code:

```
void main ()  
{  
    int i;  
    double x, pi, step, sum = 0.0;  
    long num_steps = 100000;  
  
    step = 1.0/(double) num_steps;  
  
    for (i=1;i<= num_steps; i++) {  
        x = (i-0.5)*step;  
        sum = sum + 4.0/(1.0+x*x);  
    }  
    pi = step * sum;  
}
```

Data/work distribution: each processor responsible for computing some rectangles and have private copy of `sum`. Reduce partial (local) copies of `sum` into final value

Distributed memory programming: example (cont.)

MPI code:

```
#include <mpi.h>
void main ()
{
    int i;
    double x, pi, step, sum=0.0;
    long num_steps = 100000;
    int numprocs, myid;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    step = 1.0/(double) num_steps;

    for (i = myid; i < num_steps; i += numprocs) {
        x = (i + 0.5)*step;
        sum += (4.0 / (1.0 + x*x));
    }
    mypi = step * sum;

    MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

    MPI_Finalize();
}
```

Distributed memory programming: example (cont.)

Hybrid MPI/OpenMP code:

```
#include <mpi.h>
void main ()
{
    int i;
    double x, pi, step, sum=0.0;
    long num_steps = 100000;
    int numprocs, myid;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    step = 1.0/(double) num_steps;

#pragma omp parallel for private(x) reduction(+:sum)
    for (i = myid; i < num_steps; i += numprocs) {
        x = (i + 0.5)*step;
        sum += (4.0 / (1.0 + x*x));
    }
    mypi = step * sum;

    MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

    MPI_Finalize();
}
```

Distributed memory programming: example (cont.)

Erlang code:

```
-module(pi).
-author("Zvi").
-compile([export_all]).

%% functional and parallel PI calculation
test()->
    N = 1000000,
    Impls = [serial,parallel],
    [ Impl,timer:tc(pi,calc_pi,[Impl,N]) || Impl<-Impls ].

%% 1. serial - tail recursion - increment index
calc_pi(serial,N) ->
    calc_pi_aux(1,N,1/N);

%% 2. parallel SMP version - process per scheduler
calc_pi(parallel,N) ->
    NWorkers = erlang:system_info(schedulers),
    NPerWorker = trunc(N/NWorkers+0.5),
    Step = 1/N,
    Self = self(),
    Pids = [spawn(fun() -> Self !
        sum,calc_pi_aux(I,lists:min([N,I+NPerWorker-1]),Step) end)
        || I<-lists:seq(1,N,NPerWorker)],
    lists:sum([receive sum,S -> S end || _<-Pids]).
```

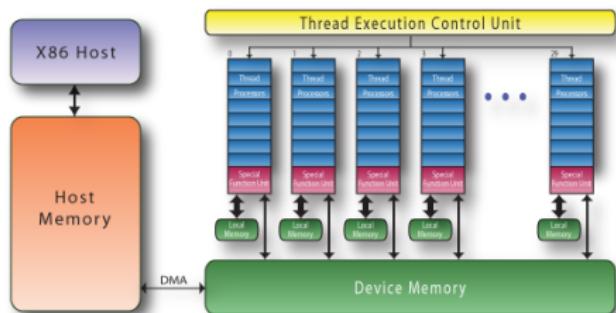
Distributed memory programming: example (cont.)

Erlang code (cont.):

```
%% Auxiliary function
calc_pi_aux(From,To,Step) ->
    calc_pi_aux(From,From,To,Step,0).

calc_pi_aux(To,_From,To,Step,Sum) -> Step*Sum;
calc_pi_aux( I,From,To,Step,Sum) ->
    X=Step*(I-0.5),
    V=4.0/(1.0+X*X),
    calc_pi_aux(I+1,From,To,Step,Sum+V).
```

GPU-enabled architectures: address space



Programmer needs

- ▶ Offload computation: low parallelism in CPU, data-parallel part in GPU
- ▶ Transfer data to devices and between devices: explicit communication
- ▶ Use synchronization mechanisms to avoid data races (between cores in GPU and between CPU-GPU)

GPU programming: example

Computation kernel for GPU:

```
#include <cuda.h>
#include <stdio.h>

// XXX float sums[size];

__global__ void pi_steps( float* sums, int size, float step, int num_steps,
    int steps_per_thread ) {

    int k = blockIdx.x * blockDim.x + threadIdx.x;
    int i = k * steps_per_thread;
    if( k < size ) {
        sums[k] = 0.0f;
        for( int j = 1; j <= steps_per_thread; j++ )
            if( i+j <= num_steps ) {
                float x = ((float)(i+j) - 0.5f) * (float)step;
                sums[k] += 4.0f/(1.0f + x*x);
            }
    }
}
```

GPU programming: example

Computation kernel for GPU (cont.):

```
// XXX: sizeof(reds) MIN size/reds_per_thread

__global__ void pi_reduction( float* reds, float* sums, int size, int reds_per_thread )
{
    int k = blockIdx.x * blockDim.x + threadIdx.x;
    int i = k * reds_per_thread;
    if( k < size/reds_per_thread ) {
        reds[k] = 0.0f;
        for( int j = 0; j < reds_per_thread; j++ )
            if( i+j < size ) reds[k] += sums[i+j];
    }
}
```

GPU programming: example

Host code:

```
#define NUM_STEPS 1048576
#define REDS_PER_THREAD 2
#define STEPS_PER_THREAD 8
#define MIN_RED 256
#define THREADS_PER_BLOCK 256

int main( ) {
    float *cu_sums, *cu_reds;
    float *red = (float*) malloc( MIN_RED * sizeof(float) );

    int num_sums = NUM_STEPS/STEPS_PER_THREAD;
    int reds_size = num_sums/REDS_PER_THREAD;
    cudaMalloc( (void**)&cu_sums, num_sums * sizeof(float) );
    cudaMalloc( (void**)&cu_reds, reds_size * sizeof(float) );
    ...
}
```

GPU programming: example

Host code (cont.):

```
...
dim3 block( THREADS_PER_BLOCK );
dim3 grid( num_sums / THREADS_PER_BLOCK );

float step = 1.0f/(float)NUM_STEPS;
pi_steps<<<grid,block>>>( cu_sums, num_sums, step, NUM_STEPS, STEPS_PER_THREAD );

float* t_sums = cu_sums;
float* t_reds = cu_reds;
int t_sums_size = num_sums;
while( t_sums_size > MIN_RED ) {
    dim3 block_red( THREADS_PER_BLOCK );
    dim3 grid_red( reds_size / THREADS_PER_BLOCK );
    pi_reduction<<<grid_red,block_red>>>( t_reds, t_sums, t_sums_size, REDS_PER_THREAD );
    t_sums_size /= REDS_PER_THREAD;
    float* tmp = t_reds;
    t_reds = t_sums;
    t_sums = tmp;
}
...
...
```

GPU programming: example

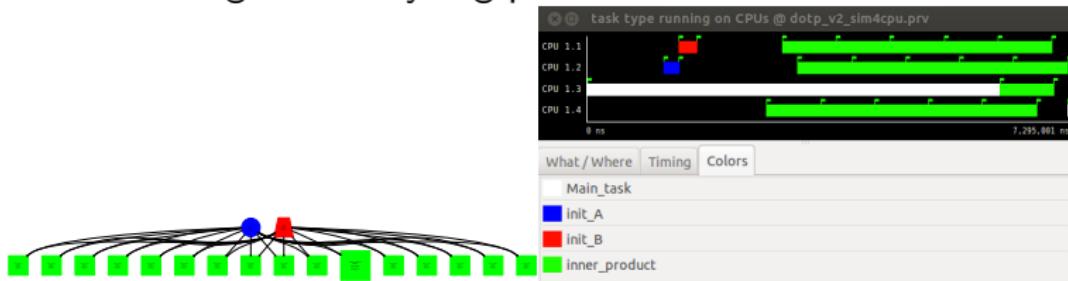
Host code (cont.):

```
...
cudaMemcpy( red, t_sums, sizeof(float) * MIN_RED, cudaMemcpyDeviceToHost );

float sum_gpu = 0.0f;
for( int i = 0; i < MIN_RED; i++ )
    sum_gpu += red[i];
float pi_gpu = step*sum_gpu;
printf("Pi GPU = %f \n", pi_gpu );
return 0;
}
```

Organization of Module II: Parallelism

- ▶ Modeling and analyzing performance



- ▶ Distributed-memory programming using MPI
- ▶ Programming GPU devices for computation acceleration using CUDA



Concurrency, Parallelism and Distributed Systems (CPDS)

Understanding parallelism

Eduard Ayguadé, Josep-Ramon Herrero, Daniel Jiménez
({eduard,josepr,djimenez}@ac.upc.edu)

Computer Architecture Department
Universitat Politècnica de Catalunya

2014/15 - Fall