


Balls and Bins: Hashing

November 3, 2019

Balls and Bins

Basic Model: Given n distinguishable (labelled) balls we throw each one **independently and uniformly** into a set of the distinguishable (labelled) bins.

$$\Pr[\text{ball } i \rightarrow \text{bin } j] = \frac{1}{m}.$$



Probability space: $\Omega = \{(b_1, b_2, \dots, b_n)\}$ where b_i denotes the index of the bin containing ball i -th. ball: $|\Omega| = m^n$.

For any $w \in \Omega$, $\Pr[w] = \left(\frac{1}{m}\right)^n$

Balls and Bins as a model

Balls and Bins as a model, is very useful in different areas of problems in computer science. For ex.:

- ▶ The **hashing data structure**: keys are the balls and the slots in the array are the bins.
- ▶ Many situations in **routing in nets**: balls represent the connectivity requirements and the bins are the paths in the network
- ▶ The **load balancing randomized algorithm**, balls are the streaming jobs and the bins are the servers.

Recall as an application of Chernoff+UB, we proved that for n balls (jobs) and m bins (servers), under a uniform and independent distribution of jobs to servers, for $n \gg m$, the probability the load of a server deviates from the expected load, was $1/m^2$.


General rules for the analysis of Balls & Bins

n balls to m bins.

- ▶ X_i a random variable counting the number of balls into bin- i .
Then $X_i \in B(n, \frac{1}{m})$.
- ▶ As we know: X_1, \dots, X_m are not independent.
- ▶ The average load in a bin is $\mu = \mathbf{E}[X_i] = n/m$.
- ▶ Rule of thumb to do the analysis:
 - ▶ If $n \gg m$, (μ large) use Chernoff bounds,
 - ▶ if $n = m$, ($\mu \in \Theta(1)$), use the Poisson approximation.

Recall that for small x ,

$$e^{-x} \sim 1 - x.$$



The Poisson Distribution

Recall that for $X \in B(n, p)$ if for large n and small p , we can have a good approximation: $\Pr[X = k] = \frac{e^{-\lambda} \lambda^k}{k!}$, where $\lambda = \mathbf{E}[X] = \mu = pn$.

For any $\lambda \in \mathbb{R}^+$, a r.v. X is said to have a Poisson $P(\lambda)$ distribution, if its PMF is $p_X(k) = \frac{e^{-\lambda} \lambda^k}{k!}$, for any $k = 0, 1, 2, 3, \dots$

Notice p_X is a correct PMF, as:

$$\sum_{k=0}^{\infty} \frac{e^{-\lambda} \lambda^k}{k!} = e^{-\lambda} (1 + \lambda + \frac{\lambda^2}{2!} + \dots) = e^{-\lambda} e^{\lambda} = 1.$$

Poisson is one of the most "natural" distributions: number of typos, number of rain drops in a square meter of roof, etc..

The Poisson Distribution: Basic Properties

Assume that $Y \in P(\lambda)$ approximates $X \in B(n, p)$, then as $\mathbf{E}[X] = np$ seems natural that $\mathbf{E}[Y] = np = \lambda$ and as $\mathbf{Var}[X] = np(1-p) = \lambda(1-p)$ and as p is small $\mathbf{Var}[X] \sim \lambda$ and $\mathbf{Var}[Y] = \lambda$. Formally, If $Y \in P(\lambda)$:

- $\mathbf{E}[Y] = \lambda$.

$$\begin{aligned}\mathbf{E}[Y] &= \sum_{k=0}^{\infty} k \frac{e^{-\lambda} \lambda^k}{k!} = e^{-\lambda} \left(\lambda + \frac{2\lambda^2}{2!} + \frac{3\lambda^2}{3!} \dots \right) \\ &= e^{-\lambda} \lambda \left(1 + \lambda + \frac{2\lambda^2}{2!} + \frac{3\lambda^2}{3!} \dots \right) = e^{-\lambda} \lambda e^{\lambda}\end{aligned}$$

Variance of Poisson r.v.

- **Var** $[Y] = \lambda$.

To prove it, instead of computing $\mathbf{E} [X^2]$ we compute $\mathbf{E} [X(X - 1)]$.

Notice **Var** $[X] = \mathbf{E} [X^2] - \mathbf{E} [X]^2 = \mathbf{E} [X(X - 1)] + \mathbf{E} [X] - \mathbf{E} [X]^2$.

$$\begin{aligned}\mathbf{E} [X(X - 1)] &= \sum_{x=0}^{\infty} x(x - 1) \frac{\lambda^x e^{-\lambda}}{x!} = \sum_{x=2}^{\infty} \frac{\lambda^2 \lambda^{x-2} e^{-\lambda}}{(x - 2)!} \\&= e^{-\lambda} \lambda^2 \sum_{x=2}^{\infty} \frac{\lambda^{x-2}}{(x - 2)!} \underbrace{=}_{y=x-2} e^{-\lambda} \lambda^2 \sum_{y=0}^{\infty} \frac{\lambda^y}{(y)!} \\&= e^{-\lambda} \lambda^2 e^{\lambda}\end{aligned}$$

So, **Var** $[X] = \lambda^2 + \lambda - \lambda^2$

Sum of Poisson r. v.

Lemma If $Y \in P(\lambda)$ and $Z \in P(\lambda')$ are independent, then $Y + Z \in P(\lambda + \lambda')$.

Proof

$$\begin{aligned}\Pr[Y + Z = j] &= \sum_{k=0}^j \Pr[(Y = k) \cap (Z = j - k)] = \sum_{k=0}^j \frac{e^{-\lambda} e^{-\lambda'} \lambda^k \lambda'^{j-k}}{k!(j-k)!} \\&= \frac{e^{-(\lambda+\lambda')}}{j!} \sum_{k=0}^j \frac{j!}{k!(j-k)!} \lambda^k \lambda'^{j-k} = \frac{e^{-(\lambda+\lambda')}}{j!} \sum_{k=0}^j \binom{j}{k} \lambda^k (\lambda')^{j-k} \\&= \frac{e^{-(\lambda+\lambda')} \times (\lambda + \lambda')^j}{j!} \Rightarrow (Y + Z) \in P(\lambda + \lambda') \quad \square\end{aligned}$$

Basic facts

Recall X_i counts the number of balls in i -th bin.

► Probability all n balls fell in the same bin: $(\frac{1}{m})^n$.

► Probability that bin i is empty:

$$\Pr[X_i = 0] = (1 - \frac{1}{m})^n \sim e^{-\frac{n}{m}} = e^{-\lambda}.$$

► Let Y be number of empty bins, compute $\mathbf{E}[Y]$?

For $1 \leq i \leq m$, let Y_i be an i.i.d. such that $Y_i = 1$ iff bin i is empty. Then,

$$\mathbf{E}[Y] = \sum_{i=1}^m \mathbf{E}[Y_i] = \sum_{i=1}^m \Pr[X_i = 0] = m(1 - 1/m)^n. \text{ So}$$

$$\mathbf{E}[Y] \sim me^{-\lambda}.$$

Probability the i -th. bin contains 1 ball

We can assume that m and n are large, (so $p = 1/m$ is small),
 $\lambda = n/m = \Theta(1)$

Exact computation: $\Pr[X_i = 1] = \binom{n}{1}(1/m)^1(1 - 1/m)^{n-1}$,
where $\binom{n}{1}$ number choices exactly 1 ball goes into bin i ,

$(1 - 1/m)^{n-1}$: remaining balls do not go to bin i .

$$\Pr[X_i = 1] = \frac{n}{m}(1 - 1/m)^n(1 - 1/m)^{-1}$$

Poisson approximation: Taking $\lambda = \frac{n}{m}$ and $(1 - 1/m)^n \sim e^{-\lambda}$ and
noticing $(1 - 1/m) \rightarrow 1$:

$$\Pr[X_i = 1] = \lambda e^{-\lambda}.$$

For $n = 3000$ and $m = 1000$, $\lambda = 3$, the exact value of
 $\Pr[X_i = 1] = 0.149286$ and the Poisson approximation is 0.149361.

Probability the i -th. bin contains exactly r balls

We can assume that m and n are large, $n, m > r$,

Exact computation: $\Pr[X_i = r] = \binom{n}{r} (1/m)^r (1 - 1/m)^{n-r}$.

Poisson approximation:

$$(1 - 1/m)^{n-r} = (1 - 1/m)^n (1 - 1/m)^{-r} = e^{-\lambda} \cdot 1^{-r}$$

$$\begin{aligned} \binom{n}{r} (1/m)^r &= \frac{1}{r!} \left(\frac{n}{m} \frac{n-1}{m} \cdots \frac{n-r+1}{m} \right) \\ &= \frac{1}{r!} \lambda \left(1 - \frac{1}{n}\right) \cdots \lambda \left(1 - \frac{r-1}{n}\right) = \lambda^r \end{aligned}$$

$$\therefore \Pr[X_i = r] \sim \frac{\lambda^r e^{-\lambda}}{r!}$$

For $n = 4000$ and $m = 2000$, $\lambda = 2$, and $r = 100$, the exact value of $\Pr[X_i = r] = 5.54572 \times 10^{-130}$ and the approximation is 1.83826×10^{-130}

Probability at least one bin has a collision

$\Pr[\text{at least 1 bin } i \text{ has } X_i > 1] = 1 - \Pr[\text{every bin } i \text{ has } X_i \leq 1]$.

If $k - 1$ balls went to $k - 1$ different bins. Then,

$\Pr[\text{The } k\text{th. ball goes into a non-empty bin}] = \frac{k-1}{m}$

$\Pr[\text{The } k\text{th. ball goes into an empty bin}] = (1 - \frac{k-1}{m})$

$$\begin{aligned}\Pr[\text{every bin } i \text{ has } X_i \leq 1] &= \prod_{i=1}^{n-1} \left(1 - \frac{i-1}{m}\right) \sim \prod_{i=1}^{n-1} e^{-i/m} \\ &= e^{-\sum_{i=1}^{n-1} i/m} = e^{-\frac{1}{m} \sum_{i=1}^{n-1} i} = e^{-\frac{n(n-1)}{2m}} \\ &\sim e^{-\frac{n^2}{2m}}\end{aligned}$$

Therefore, $\Pr[\text{at least 1 bin } i \text{ has } X_i > 1] \sim 1 - e^{-\frac{n^2}{2m}}$.

Birthday problem

How many students in a class, to have that with probability $> 1/2$ at least 2 have the same birthday

This is the same problem as above, with $m = 365$:

$$\begin{aligned}\text{We need } e^{-\frac{n^2}{2m}} &\leq \frac{1}{2} \Rightarrow \frac{n^2}{2m} \leq \ln 2 \\ \Rightarrow n &= \sqrt{2m \ln 2}. \text{ If } m = 365 \text{ then } n = 22.49.\end{aligned}$$

Therefore, If there are more than 23 students in a class, with probability greater than $1/2$, more than 2 students will have the same birthday

Coupon Collector's problem

A. de Moivre (VIIc.)

How many balls do we need to throw to assure that w.h.p. every bin contains ≥ 1 balls

- ▶ Let Y a r.v. counting the number of balls we have to throw until having no empty bins
- ▶ For $i \in [m]$, let $Y_i = \#$ balls between between $i - 1$ bins are not empty and the bin i gets a ball.
- ▶ So $Y = \sum_{i=1}^m Y_i$
- ▶ For $i \in [m]$, define Z_i a r.v. $\#$ balls until first ball goes into $\rightarrow i$ -bin.
- ▶ Then $Y = Z_m$, and for $i \in [m]$, let $Y_1 = Z_1$, for $i \geq 2$, $Y_i = Z_i - Z_{i-1}$
- ▶ First we want $\mathbf{E}[Y] = \sum_{i=1}^n \mathbf{E}[Y_i]$. After we prove concentration

Coupon Collector's problem

$Y_i = \#$ of balls we have to throw to get a new non-empty bin (it will be the i -th. non-empty bin)

$\Pr[\text{a new ball going into non-empty bin}] = 1 - \frac{1}{m}$.

If $k = \#$ balls between $(i-1)$ and i :

$$\Pr[Y_i = k] = \left(\frac{i-1}{m}\right)^{k-1} \underbrace{\left(1 - \frac{i-1}{m}\right)}_{p_i}.$$

Therefore $Y_i \in G(p_i)$ and $\mathbf{E}[Y_i] = \frac{m}{m-i+1}$.

$$\mathbf{E}[Y] = \sum_{i=1}^m \mathbf{E}[Y_i] = \sum_{i=1}^n \frac{m}{m-i+1} = n \sum_{j=1}^n \frac{1}{j} = n(\ln n + o(1)).$$

Coupon Collector's problem: Concentration

Let $\mathbf{E}[Y] = O(\ln m) \sim cm \ln m$ for constant $c > 1$

- ▶ For any ball i , define the event A_j^r : bin $j = \emptyset$ after the first r throws.
- ▶ Notice events $A_1^r, A_2^r, \dots, A_m^r$ are not independent.
- ▶ $\Pr[A_j^r] = (1 - \frac{1}{m})^r \sim e^{-r/n} \leq e^{-cm \ln m / m} = n^{-c}$.
- ▶ Let W be a r.v. counting the number of balls needed so every bin has load ≥ 1 .

$$\begin{aligned}\Pr[W > cm \lg m] &= \Pr\left[\bigcup_{j=1}^m A_j^{cm \lg m}\right] \underbrace{\leq}_{UB} \sum_{j=1}^m \Pr[A_j^{cm \lg m}] \\ &\leq \sum_{j=1}^m n^{-c} = m^{1-c}.\end{aligned}$$

$$\Pr[W > cm \lg m] \leq n^{1-c}.$$

Coupon Collector's problem: Concentration Bounds

- ▶ The previous bound using UB is more tight than the one using Chebyshev or Chernoff on random variable Y .
(See homework)
- ▶ In Section 5.4.1 of MU book, there is a sharper bound for the Coupon collector's, using the Poisson approximation.

Maximum Load

This is a very similar problem to the job and servers, but with sharper bounds

Theorem If we throw n balls independently and uniformly into $m = n$ bins, then the maximum loaded of a bin as at most $\left(\frac{4 \lg n}{\lg \lg n}\right)$ balls, with probability $\leq 1 - \frac{1}{n}$, i.e. w.h.p.

Recall that if for any bin $1 \leq j \leq n$, $X_j =$ is a r.v. with its load.

We know $\{X_j\}$ are not independent and $\mathbf{E}[X_j] = n/n = 1$.

To show the above bound we use the following two inequalities:

$$\left(\frac{N}{K}\right)^K \leq \binom{N}{K} \leq \left(\frac{Ne}{K}\right)^K. \quad (1)$$

$$\text{Let } N > e. \text{ If } K \geq \frac{2 \ln N}{\ln \ln N} \text{ then } K^K \geq N. \quad (2)$$

Max-load: Proof Upper Bound

For $1 \leq k \leq n$, $\Pr[X_i \geq k] \leq \binom{n}{k} \frac{1}{n^k} \leq \left(\frac{ne}{k}\right)^k \frac{1}{n^k} \leq \left(\frac{e}{k}\right)^k$.

We want to prove that for $k \geq \frac{2 \ln n}{\ln \ln n} \Rightarrow \Pr[X_i \geq \frac{2 \ln n}{\ln \ln n}] \leq \frac{1}{n^2}$.

i.e. $\Pr[X_i \geq k] \leq \left(\frac{e}{k}\right)^k \leq \frac{1}{n^2} \Rightarrow \left(\frac{e}{k}\right)^{\frac{k}{e}} \geq n^{\frac{2}{e}}$

Taking \ln : $\frac{k}{e} \geq \frac{2 \ln(n^{2/e})}{\ln \ln(n^{2/e})} = \frac{4 \ln n}{e \ln(\frac{2}{e} \ln n)} \Rightarrow k \geq \frac{4 \ln n}{\ln(\frac{2}{e} \ln n)}$

We proved that if $k \geq \frac{4 \ln(n)}{\ln(2/e) \ln \ln(n)}$ then $\Pr[X_i \geq k] \leq \frac{1}{n^2}$.

Then, using U-B

$\Pr[\exists i \in [n] \mid X_i \geq k] \leq \sum_{i=1}^n \Pr[X_i \geq k] \leq \frac{n}{n^2} = \frac{1}{n}$.

Further considerations on Max-load

1. The same proof could be extended to the case of n balls and m bins, with the constrain $n < m \ln m$.
2. We can obtain the same result by using Chernoff's bounds.
(Nice exercise!)
3. In fact, the result could be extended to prove the Lower Bound: that w.h.p. the max-load is $\Omega(\frac{\ln n}{\ln \ln(n)})$ balls. One easy way to prove the lower bound is using Chebyshev's bound.
4. That result yields: Throwing n balls to n bins, w.h.p. we have a max-load of $\Theta(\frac{\ln n}{\ln \ln(n)})$.
5. We can obtain sharper bounds for max-load, using strong inequalities (Azuma-Hoeffding) or the Poisson approximation.

Poisson approximation

1. A difficulty with the **exact** (binomial) B & B model is that random variables could be dependent (for ex. bin's load).
2. We have seen how to approximate the expressions arising from the exact computations by a Poisson, **if p is small and n is large**.
3. However, under the right conditions, we can approach the whole solution to the problem by using Poisson r.v. instead of Binomial. In the binomial case we have exactly n balls with probability $p = 1/m$, in the Poisson case we have an intensity $\lambda = n/m$, where n is the expected number of balls being used.
4. The Poisson case is to use directly independent Poisson random variables and it can be shown, under certain conditions give a good approximation to the solution. See for ex. section 5.4 in MU.

Dynamic Sets.

Given a **universe** \mathcal{U} and a set of **keys** $\mathcal{S} \subset \mathcal{U}$, for any $k \in \mathcal{S}$ we can consider the following operations

- ▶ **Search** (\mathcal{S}, k) : decide if $k \in \mathcal{S}$
- ▶ **Insert** (\mathcal{S}, k) : $\mathcal{S} := \mathcal{S} \cup \{k\}$
- ▶ **Delete** (\mathcal{S}, k) : $\mathcal{S} := \mathcal{S} \setminus \{k\}$
- ▶ **Minimum** (\mathcal{S}) : Returns element of \mathcal{S} with smallest k
- ▶ **Maximum** (\mathcal{S}) : Returns element of \mathcal{S} with largest k
- ▶ **Successor** (\mathcal{S}, k) : Returns element of \mathcal{S} with next larger key to k
- ▶ **Predecessor** (\mathcal{S}, k) : Returns element of \mathcal{S} with next smaller key to k .

Recall Dynamic Data Structures

DICTIONARY

Data structure for maintaining $\mathcal{S} \subset \mathcal{U}$ together with operations:

- ▶ **Search** (\mathcal{S}, k) : decide if $k \in \mathcal{S}$
- ▶ **Insert** (\mathcal{S}, k) : $\mathcal{S} := \mathcal{S} \cup \{k\}$
- ▶ **Delete** (\mathcal{S}, k) : $\mathcal{S} := \mathcal{S} \setminus \{k\}$

PRIORITY QUEUE

Data structure for maintaining $\mathcal{S} \subset \mathcal{U}$ together with operations:

- ▶ **Insert** (\mathcal{S}, k) : $\mathcal{S} := \mathcal{S} \cup \{k\}$
- ▶ **Maximum** (\mathcal{S}) : Returns element of \mathcal{S} with largest k
- ▶ **Extract-Maximum** (\mathcal{S}) : Returns and erase from \mathcal{S} the element of \mathcal{S} with largest k

Hashing functions


Data Structure that supports **dictionary** operations on an universe of **numerical** keys.

Notice the number of possible keys represented as 64-bit integers is $2^{64} = 18446744073709551616$.

Tradeoff **time/space**

Define a **hashing table** $T[0, \dots, m-1]$

a **hashing function** $h : \mathcal{U} \rightarrow T[0, \dots, m-1]$



Simple uniform hashing function.

A good hashing function must have the property that $\forall k \in \mathcal{U}$, $h(k)$ must have the **same probability** of ending in any $T[i]$.

Given a hashing table T with m slots, we want to store $n = |\mathcal{S}|$ keys, as maximum.

Important measure: **load factor** $\alpha = n/m$, the average number of keys per slot.

The performance of hashing depends on how well h distributes the keys on the m slots: h is **simple uniform** if it hash any key *with equal probability* into any slot, independently of where other keys go.

How to choose h : The division method

Choose m prime and as far as possible from a power,

$$h(k) = k \bmod m.$$

Fast ($\Theta(1)$) to compute in most languages ($k \% m$)!

Be aware: if $m = 2^r$ the hash does not depend on all the bits of K


If $r = 6$ with $k = 1011000111 \underbrace{011010}_{=h(k)}$

$(45530 \bmod 64 = 858 \bmod 64)$

Collision resolution: Separate chaining

For each table address, construct a linked list of the items whose keys hash to that address.

- ▶ Every key goes to the same slot
- ▶ Time to explore the list = length of the list



Cost of exploring the list

The cost of the dictionary operations:

- ▶ Insertion of a new key: $\Theta(1)$.
- ▶ Search of a key: $O(\text{length of the list})$
- ▶ Deletion of a key: $O(\text{length of the list})$.

Under the hypothesis that h is *simply uniform*, the expected number of keys falling into $T[i]$ is $\alpha = n/m$.

Therefore, the expected time to search the list at $T[i]$ is $O(1 + \alpha)$.

Theorem

Under the assumption of simple uniform hashing, in a hash table with chaining, an unsuccessful and successful search takes time $\Theta(1 + \frac{n}{m})$ on the average.

Bloom filter

Given a set of elements S , we want a Data structure for supporting insertions and querying about membership in S .

In particular we wish a DS s.t.

- ▶ *minimizes* the use of memory,
- ▶ *can check membership as fast* as possible.

Burton Bloom: The Bloom filter data structure. Comm. ACM, July 1970.

A hash data structure where each register in the table is one bit

Definition Bloom filter

Create a **one bit** hash table $T[0, \dots, m-1]$, and a hash function h . Initially all m bits are set to 0.

Given a set $S = \{x_1, \dots, x_n\}$ define a hashing function $h : S \rightarrow T$. For every $x_i \in S$, $h(x_i) \rightarrow T[j]$ and $T[j] := 1$.


Given a set S a function $h()$ and a table $T[m]$:

```
Insert ( $x$ )  
 $h(x) \rightarrow i$   
if  $T[i] == 0$  then  
     $T[i] = 1$   
end if
```

```
inS( $y$ )  
 $h(y) \rightarrow i$   
if  $T[i] == 1$  then  
    return Yes  
else  
    return No  
end if
```

Notice: once we have hashed S into T we can **erase** S .

False positives



Bloom filter needs $O(m)$ space and answers membership queries in $\Theta(1)$.

Inconvenience: Do not support removal and may have **false positive**.

In a query $y \in S?$, a Bloom filter always will report correctly if indeed $y \in S$ ($h(y) \rightarrow T[i]$ with $T[i] = 1$), but if $y \notin S$ it may be the case that $h(y) \rightarrow T[i]$ with $T[i] = 1$, which is called a **False positive**.

How large is the error of having a false positive?

Probability of having a false positives

Let $|S| = n$, we constructed a BF $(h, T[m])$ with all elements in S . If we query about $y \in S?$, with $y \notin S$, and $h(y) \rightarrow T[i]$, what is the probability that $T[i] = 1$?

After all the elements of S are hashed into the Bloom filter, the probability that a specific $T[i] = 0$ after hashing n elements is $(1 - \frac{1}{m})^n = e^{-n/m}$

Therefore, for a $y \notin S$, the probability of false positive π :

$$\pi = \Pr[h(y) \rightarrow T[i] \mid T[i] = 1] = 1 - (1 - \frac{1}{m})^n \sim 1 - e^{-n/m}.$$


To minimize π , $1 - e^{-n/m}$ has to be small, $\Rightarrow 1/e^{n/m}$ small, i.e., $m \gg n$.

For ex.: if $m = 100n$, $\pi = 0.0095$; If $m = n$, $\pi = 0.632$ and if $m = n/10$, $\pi = 0.9999$

Alternative: Amplify

Take k different functions $\{h_1, h_2, \dots, h_k\}$ in the same 2-universal set of functions.

Ex. Bloom filter with 3 hash functions: h_1, h_2, h_3 .



When making a query about if $y \in S$, compute $h_1(y), \dots, h_t(y)$, if one of them is 0 we certainly $y \notin S$, else (if all the k hashing go to bits with value 1) $y \in S$ with some probability.

After hashing the n elements k times to T , for an specific $T[i]$:

$$p = \Pr[T[i] = 0] = \left(1 - \frac{1}{m}\right)^{kn} = e^{-kn/m}.$$

The probability f of a false positive:

$$f = \left(1 - e^{-kn/m}\right)^k = (1 - p)^k$$

Optimizing k

Given n and m we want to find the optimal value of k to minimize the probability of a false positive $f(k) = (1 - e^{-kn/m})^k$

Define $g(k) = \ln f(k) = k \ln(1 - e^{-kn/m})$. Minimizing f is equivalent to minimizing g .

To minimize the probability of having a false positive: $\frac{dg(k)}{dk} = 0$

$$\Rightarrow \frac{dg(k)}{dk} = \ln(1 - e^{-kn/m}) + \frac{kne^{-kn/m}}{m(1 - e^{-kn/m})} = 0,$$

\Rightarrow when n, m are given, to minimize f is $k_o = (\ln 2) \frac{m}{n}$.

In this case the false positive probability $f_o = 0.6185^{m/n}$.

Bloom filters allow a constant probability of false positive, $m = cn$ for small constant c , i.e. m grows linear wrt n .

For ex.: if $c = 2$ and $k = 6$ the false positive probability is around 2%.

Practical issues

For password checking:

If D has 100000 common words, each of 7 characters \Rightarrow we need 700000 bytes

Use 5 tables of 160000 bits each \Rightarrow need a total of 800000 bits = 100000 bytes.

The probability of error is 0.02

On the other hand although the results shown before are asymptotic, there also work for practical values of n . Figure in the side table give the probability of false positive wrt to n

